

UNIVERSITY *of* WASHINGTON

Week 3: Intro to SQL



W

AGENDA

- > Homework Overview
- > MariaDB
- > Introduction to SQL
 - Creating/Altering/Dropping Tables
 - Inserting Tuples
 - Selection, Joins





HOMEWORK 1

Overall good job!

I will have grades posted this weekend.

Let's talk about solutions.

W

UNIVERSITY *of* WASHINGTON

Introduction to RDBMs: MariaDB



Which RDBMS will we use?

- > MariaDB: Drop-in replacement of MySQL
 - Open Source
 - “Compatible” with MySQL GUIs such as MySQL Workbench, Navicat, Toad, etc.
- > “Created” when Oracle acquired MySQL


<https://mariadb.com/kb/en/mariadb/documentation/>
Google: “Maria DB Documentation”




Client/Server vs Local



Entirely Local



SSH Connection



Hybrid (Hosted)



MariaDB Binaries



MariaDB Database



Entirely Local

- > Install MariaDB
- > Configure your root level access
- > Act as your own DBA, create schemas, databases etc, play with DBMS level settings
- > To connect: localhost (127.0.0.1)
- > Make sure your DB Server is running.



Run entirely from cssgate

- > **cssgate.insttech.washington.edu runs a version of MariaDB and allows for remote connections**
- > **Can control via Terminal or PUTTY vis SSH**
- > **ssh uwnetid@cssgate.insttech.washington.edu**
- > **Or connect via GUI (explained later)**



Hybrid

- > Install MariaDB
- > Connect to cssgate as a remote server

```
mysql -u uwnet -h cssgate.insttech.washington.edu -p
```

- > Use your sql password (obtained from ~/.pw file)



MariaDB GUI / IDE

- > Many options for your IDE
 - TOAD, MySQL Workbench, NaviCat to name a few...
- > We are going to use JetBrains: IntelliJ iDEA
 - Free for Students



If you don't already have a student license...



<https://www.jetbrains.com/student/>

Apply for License

Confirm Email Address

Activate Account

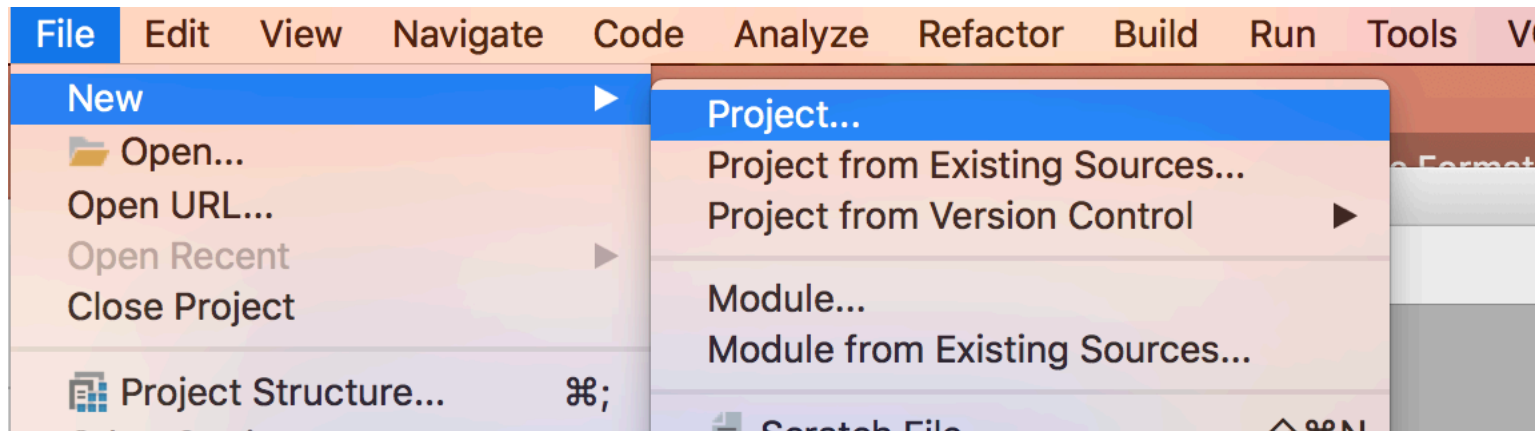
Create Username

Download and Install IntelliJ IDEA Ultimate



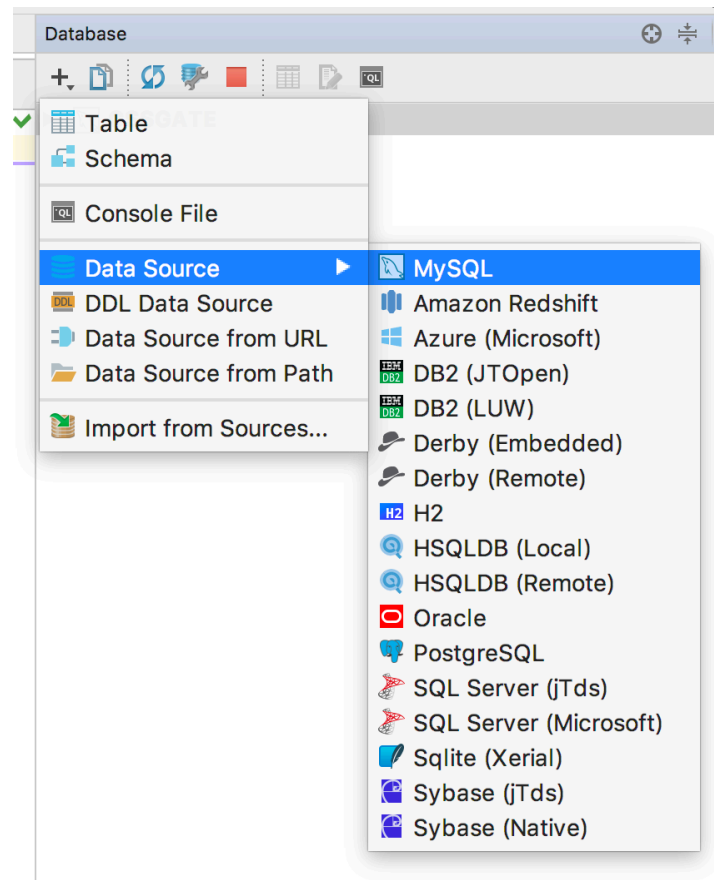
Configuring your GUI

- > Create a new Project for TCSS445
 - If not prompted:

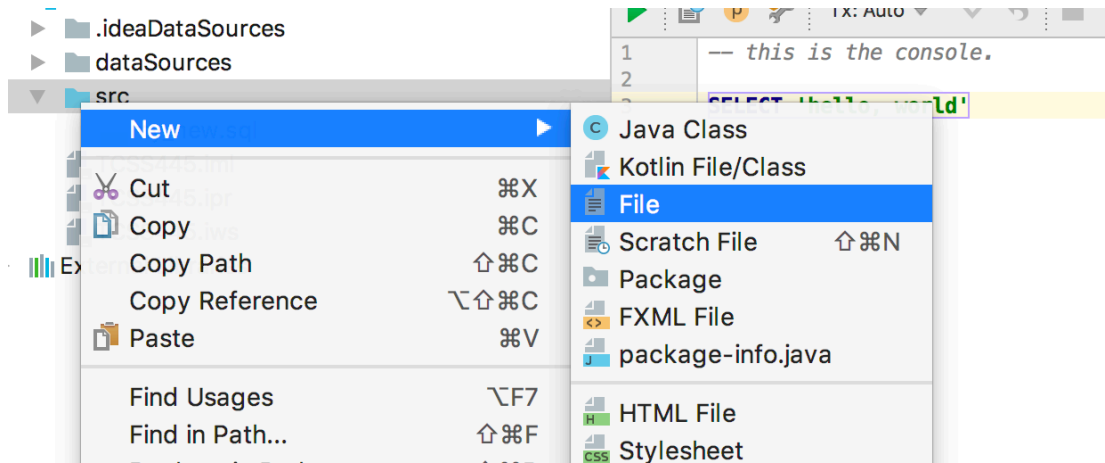


Configuring your GUI

> Tool Window > Database

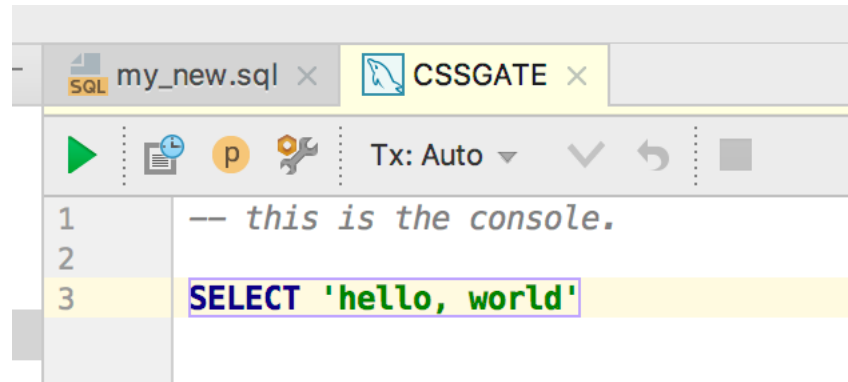


Configuring your GUI



Save queries/scripts as .sql files.

Interact without saving files
using the console...



W



UNIVERSITY *of* WASHINGTON

INTRO TO SQL



Introduction to SQL

- > Based on Relational Algebra
- > “SQL” rooted from a ‘standard’ (current SQL-11)
- > SQL standard also referred to as ANSI-SQL
- > SQL Server, Oracle, MariaDB/MySQL, all are different dialects of ANSI SQL



Types of SQL Statements

- > DDL (Data Definition Language)
 - “Create, Alter, Drop”
 - Defines the schema and functional dependencies
- > DML (Data Manipulation Language)
 - “Select, Insert, Update, Delete”
 - Retrieves manipulates data within constraints of the schema. DML cannot function outside of what is defined by DDL.
- > DCL (Data Control Language) *not in your book*
 - “Grant, Revoke”
 - Provides (or denies) users rights to data



Relations, Tables, Views, Temp

- > RDBMSs implement **relations** in several ways:
 - **Tables**: Persisted data. Arranged in a set way on a disk.
 - **Views**: Persisted “Queries.” Relation is created when requested.
 - **Temp Tables**: Cached Data (in memory). Can be stored globally, per session, or per query (behavior varies between RDBMS).



Creating Tables

```
CREATE TABLE <name> (  
    <colname> <datatype> <options>  
    ,<colname> <datatype> <options>  
    ,...  
    ,PRIMARY KEY(<colname>,<colname>)  
);
```

For complete syntax reference for MariaDB:
<https://mariadb.com/kb/en/library/create-table/>



Tips for Table Naming

> Choose a convention and stick to it!

- Most common casings:

- > MACRO_CASE
- > PascalCase
- > snake_case

- Rarely:

- > kebab-case
- > camelCase

- Unless you're using MS Access avoid Hungarian Notation!

- > tblData

"The wonderful thing about standards is that there are so many of them."
-Grace Murray Hopper

> Avoid Abbreviations unless completely obvious! (Unless using Oracle)

- Employee (not Emp)
- Transactions (not Txs)
- DATE_DIM (DIM for dimension is common/acceptable)

> Singular Name, not Plural

- It's assumed that there will be multiple Employees in the Employee table.



More on Singular Vs Plural

- > In your text: Movies and MovieStar are both tables in the same schema.
- > If you knew you had a movie table and a movie star table, you might try:
select * from Movie; or
select * from MovieStars
- > Some notable exceptions: People is a common table (people is plural, it isn't the Person table usually)



Back to Creating Tables...

```
CREATE TABLE Employee (  
    <colname> <datatype> <options>  
    ,<colname> <datatype> <options>  
    ,...  
    ,PRIMARY KEY(<colname>,<colname>,...)  
);
```



General Tips for Column Names

- > **Stick to a convention!**
 - TableNameId is a common Primary Key as is “ID”
- > **Common Casings:**
 - PascalCase
 - MACRO_CASE
 - snake_case
- > **Be specific, Abbreviations are more acceptable than in Table names, but avoid if possible.**
- > **Avoid Reserved words like “SELECT” and “KEY”**



Column Names, how descriptive?

> Which is “better?”

- Transaction.processing_time
- Transaction.processing_time_ms

> Which is “better?”

- Product.height
- Product.height_inches

> Alternatives?



Column Names, how descriptive?

- Transaction.processing_time
- Transaction.processing_time_unit
- Product.height
- Product.height_unit



Back to Creating Tables...

```
CREATE TABLE Employee (  
    EmployeeId <datatype> <options>  
    ,<colname> <datatype> <options>  
    ,...  
    ,PRIMARY KEY(<colname>,<colname>,...)  
);
```



Common Data Types

TINYINT (-128 to 127)	1 byte
BOOLEAN (0 & 1)	1 byte
INT (-2 ³¹ to 2 ³¹ -1)	4 bytes
FLOAT(M,D) M = total digits, D = digits after the decimal point	M <= 25 ? (4bytes) : (8bytes)
CHAR(M)	M * bytes in char set
VARCHAR(M) M < 2 ¹⁶ - 1	M * bytes per char + 2 bytes
TIME (microsecond precision)	3 bytes
DATE (1000-01-01 to 9999-12-31)	3 bytes
DATETIME (1000-01-01 to 9999-12-31)	8 bytes

Full list: <https://mariadb.com/kb/en/mariadb/data-types/>



DataTypes

> Size matters

- Storage is cheap? Why does it matter?

> Usage matters

- How is this attribute going to be used?
- 20170121 (3bytes MED INT)
- 2017-01-21 12:00: AM (4bytes, datetime)



SOME OF OUR OPTIONS

```
CREATE TABLE Employee (  
    EmployeeId INT <options>  
    ,FirstName VARCHAR(50) <options>  
    ,PRIMARY KEY(<colname>,<colname>...)  
);
```



DEFAULT VALUES

- > Make the database do some of the heavy lifting
- > Can speed up INSERT statements

```
CREATE TABLE Employee (  
    EmployeeId INT  
    ,FirstName VARCHAR(50) DEFAULT 'Unknown'  
    ,RecordCreateDate TIMESTAMP  
        DEFAULT CURRENT_TIMESTAMP  
    ,PRIMARY KEY(<colname>,<colname>,...)  
);
```



NULL VS NOT NULL

```
CREATE TABLE Employee (  
    EmployeeId INT  
    ,FirstName VARCHAR(50) NOT NULL  
    ,RecordCreateDate TIMESTAMP  
        DEFAULT CURRENT_TIMESTAMP  
    ,PRIMARY KEY(<colname>,<colname>,...)  
);
```



Primary Keys

```
CREATE TABLE Employee (  
    EmployeeId INT  
    ,FirstName VARCHAR(50) NOT NULL  
    ,RecordCreateDate TIMESTAMP  
        DEFAULT CURRENT_TIMESTAMP  
    ,PRIMARY KEY(<colname>,<colname>,...)  
);
```



PRIMARY KEYS

- > Must be unique!
- > Cannot be NULL
- > Often how the table is sorted on disk
- > Smaller the better
- > “Natural vs Generated | Artificial Keys?”
- > CHAR | VARCHAR | INT | BIGINT?!
- > AUTO_INCREMENT





UNIQUE

- > **UNIQUES MUST BE UNIQUE**
- > **BUT CAN BE NULL**



CREATED TABLE

```
CREATE TABLE Employee (  
    EmployeeId INT  
    ,FirstName VARCHAR(50) NOT NULL  
    ,RecordCreateDate TIMESTAMP  
        DEFAULT CURRENT_TIMESTAMP  
    ,PRIMARY KEY(EmployeeId)  
);
```



CREATED TABLE

```
CREATE TABLE Employee (  
    EmployeeId INT PRIMARY KEY AUTO_INCREMENT  
    ,FirstName VARCHAR(50) NOT NULL  
    ,RecordCreateDate TIMESTAMP  
        DEFAULT CURRENT_TIMESTAMP  
);
```



DROPPING A TABLE

> DROP TABLE <tablename>;



ALTERING A TABLE

```
ALTER TABLE <tablename>  
  ADD col_name col_def  
    FIRST | AFTER col_name
```

```
ALTER TABLE Employee  
  ADD LastName VARCHAR(50) DEFAULT 'unknown'  
    AFTER FirstName
```

<https://mariadb.com/kb/en/mariadb/alter-table/>



INSERTING RECORDS (BASIC)

```
INSERT INTO <TABLENAME>(<COLNAME>,...) VALUES  
    (TUPLE1, ...)  
, (TUPLE2, ...)
```

```
INSERT INTO Employee VALUES  
    (1, 'Jane', 'Doe', CURRENT_TIMESTAMP)  
, (2, 'John', 'Doe', CURRENT_TIMESTAMP)
```

```
INSERT INTO Employee (EmployeeId, FirstName, LastName) VALUES  
    (1, 'Jane', 'Doe')  
, (2, 'John', 'Doe')
```



INSERTING WITH DEFAULTS AND AUTO_INCREMENT

```
INSERT INTO Employee (FirstName, LastName) VALUES
  ('Homer','Simpson')
, ('Marge','Simpson')
, ('Maggie','Simpson')
, ('Chief','Wiggum')
```

```
SELECT * FROM Employee
```

EmployeeId	FirstName	LastName	RecordCreateDate
1	Homer	Simpson	2017-01-20 22:10:00 AM
2	Marge	Simpson	2017-01-20 22:10:00 AM
3	Maggie	Simpson	2017-01-20 22:10:00 AM
4	Chief	Wiggum	2017-01-20 22:10:00 AM





UNIVERSITY *of* WASHINGTON

DML: SELECTION



INTRO TO SELECTION

SELECT	[COLUMNS & EXPRESSIONS]
FROM	[TABLE VIEW]
JOIN	[TABLE VIEW]
ON	[CONDITIONS]
WHERE	[CONDITIONS]
GROUP BY	[COLUMNS & EXPRESSIONS]
HAVING	[CONDITIONS]
ORDER BY	[COLUMNS & EXPRESSIONS]



INTRO TO SELECTION

SELECT	[COLUMNS]	π
FROM	[TABLE]	(Relation)
JOIN	[TABLE]	\bowtie
WHERE	[CONDITIONS]	σ
GROUP BY	[COLUMNS]	γ
HAVING	[CONDITIONS]	σ
ORDER BY	[COLUMNS]	τ



SELECT (Projection)

```
SELECT <COLUMNS|EXPRESSIONS>
```

```
SELECT CURRENT_TIMESTAMP
```

```
SELECT 10*100.0
```

```
SELECT *
```



Case Sensitivity

- > In your book, it states “SQL is Case Insensitive”
 - Not true in all implementations...
- > Case sensitive file systems can result in case sensitive objects and values.
- > The keywords are NOT generally case-sensitive:
`CREATE TABLE Employee != CREATE TABLE EMPLOYEE`
`creATE Table employee = CREATE TABLE employee`



Dealing with Case Sensitivity

- > If you are using a case sensitive Database choose:
 - `MACRO_CASE`
 - `snake_case`
- > Even if not, try and keep your queries with matching case
 - I got burned on this once.



FROM (The relation(s) to be queried)

```
SELECT c1  
FROM t1
```

Is the same as:

$\pi_{c_1}(t_1)$



WHERE (Condition/Criteria)

```
SELECT c1  
FROM t1  
WHERE c1 >= 10
```

Same as:

$$\pi_{c1}(\sigma_{c1 \geq 10}(t1))$$


Exercise: Write DDL to create the following Tables

Employee (EmpId, FName, Lname, Gender, MgrId, LocId, EmailAddress)

Location (LocId, LocName, City, State)

Class (ClassId, Cname, InstructorId, LocId, TotHours)

```
CREATE TABLE <name> (  
    <colname> <datatype> <options>  
    ,<colname> <datatype> <options>  
    ,...  
    ,PRIMARY KEY(<colname>,<colname>)  
);
```



Exercise: Insert Tuples from HW1

> E.g. Employee

- 106, Petrina, Tillman, F, NULL, 2, 106@company.com
- 112, Alec, Wilhoit, M, 106, 1, 112@company.com

> E.g. Location

- 1, CoffeeTree, Seattle, WA
- 2, Evergreen, Tacoma, WA

> E.g. Class

- 1003, Conflict Management, 112, 1, 10
- 1005, Management Essentials, 112, 1, 40



Write simple select statements:

- > Get first and last names of Female Employees
- > Get location name, city and state for Locations in Washington
- > Get class names where total hours are more than 20



WHERE: PATTERN MATCHING WITH 'LIKE'

% = Any number of characters (including zero)

_ = Any Single Character

Ex: Cname LIKE 'Management%'

Ex: Urgency LIKE 'P_ Resolution'



WHERE: PATTERN MATCHING WITH 'LIKE'

> Find courses that start with the letter "P"

Find courses that are at least 3 words?



WHERE: PATTERN MATCHING WITH 'LIKE'

> Find courses that start with the letter "P"

```
WHERE Cname LIKE 'P%'
```

Find courses that are at least 3 words?

```
WHERE Cname LIKE '% % %'
```



MARIA DB SUPPORTS REGULAR EXPRESSIONS

Look up all classes that end in 1xx

(could use LIKE '%1__' but that would match 'Some class 1ab')

```
WHERE CName REGEXP '^.*1[0-9]{2}$';
```

(translated Starts with any character, any number of times, followed by a 1, directly followed by exactly 2 numeric characters, followed directly by the end of the string)





UNIVERSITY *of* WASHINGTON

JOINS!



W

Implicit Joins (ew)

```
SELECT *  
FROM Class, Location
```

Returns all columns from both tables, matching on their common attributes (Natural Join)



Implicit Joins (ew)

```
SELECT name
```

```
FROM Movies, MovieExec
```

```
WHERE title = 'Star Wars'  
      AND ProducerC# = cert#
```



Explicit Joins: KA preferred.

```
SELECT MovieExec.name
```

```
FROM Movies
```

```
INNER JOIN MovieExec
```

```
    ON Movies.ProducerC# = MovieExec.cert#
```

```
WHERE title = 'Star Wars'
```



Aliases and Disambiguation

> Columns and Tables can (and should!) be aliased

```
SELECT
  C.Cname    as ClassName
, L.LocName as LocationName
, L.City

FROM Class C
JOIN Location L
  on C.LocId = L.LocId
```

ClassName	LocationName	City
Customer Service 101	Mile High	Denver
Conflict Management	CoffeeTree	Seattle




Aliases and Disambiguation

> Remember our self join in Homework 1?

```
SELECT
    E.Fname      as EmpFirstName
  , E.Lname      as EmpLastName
  , MGR.Fname    as MgrFirstName
  , MGR.Lname    as MgrLastName

FROM Employee E
JOIN Employee MGR
    on E.MgrId = MGR.EmpId
```





UNIVERSITY *of* WASHINGTON

Unions, Intersections, and Difference



W

Set Operations

- SQL provides corresponding operators that apply to the results of queries, provided those queries produce relations with the same list of attributes and attribute types.
- The keywords used are UNION, INTERSECT, and EXCEPT for \cup , \cap , and $-$, respectively.
- There is a bag union called “UNION ALL” that will not remove duplicates but rather stick one relation on top of the other.
- *Not all databases support these keywords.*



Set Operations

MovieStar (name, address, gender, birthdate)

MovieExec (name, address, cert#, netWorth)

Movies (title, year, length, genre, studioName, producerC#)

StarsIn (movieTitle, movieYear, starName)

```
SELECT name, address
  FROM MovieStar
 WHERE gender = 'F'
INTERSECT
SELECT name, address
  FROM MovieExec
 WHERE netWorth > 10000000;
```

```
SELECT name, address
  FROM MovieStar
EXCEPT
SELECT name, address
  FROM MovieExec;
```

```
SELECT title, year
  FROM Movies
UNION
SELECT movieTitle AS title
      , movieYear  AS year
  FROM StarsIn;
```

