

UNIVERSITY *of* WASHINGTON

Week 4: SQL II: The Sequel



AGENDA

- > Quiz Reminder
- > Introduce Project
- > SQL I Review & Examples
- > More SQL
 - Subqueries
 - Relation Operators
 - Outer Joins
 - Grouping & Aggregation



SQL I REVIEW EXERCISES...

```
MOVIE (  
  title:      string,  
  year_made:   integer,  
  length_mins: integer,  
  genre:      string,  
  studio_name: string,  
  producer_num: integer  
)
```

```
MOVIE_STAR (  
  name:      string,  
  address:   string,  
  gender:    char,  
  birthdate: date  
)
```

```
STARS_IN (  
  title:      string,  
  year_made: integer,  
  star_name: string  
)
```

```
MOVIE_EXECUTIVE (  
  full_name: string,  
  address:   string,  
  cert_num: integer,  
  net_worth: integer  
)
```

```
STUDIO (  
  legal_name: string,  
  address:    string,  
  prez_cert_num: integer  
)
```



SQL I REVIEW EXERCISES...

```
DROP TABLE IF EXISTS MOVIE;  
CREATE TABLE MOVIE (  
  title VARCHAR(50)  
  , year_made INT  
  , length_mins INT  
  , genre VARCHAR(20)  
  , studio_name VARCHAR(50)  
  , producer_num INT  
  , PRIMARY KEY(title, year_made)  
);
```

```
MOVIE_STAR (  
  name:      string,  
  address:   string,  
  gender:    char,  
  birthdate: date  
)
```

```
STARS_IN (  
  title:      string,  
  year_made: integer,  
  star_name: string  
)
```

```
MOVIE_EXECUTIVE (  
  full_name: string,  
  address:   string,  
  cert_num:  integer,  
  net_worth: integer  
)
```

```
STUDIO (  
  legal_name:  string,  
  address:     string,  
  prez_cert_num: integer  
)
```



SQL I REVIEW EXERCISES...

```
DROP TABLE IF EXISTS MOVIE;  
CREATE TABLE MOVIE (  
  title VARCHAR(50)  
  , year_made INT  
  , length_mins INT  
  , genre VARCHAR(20)  
  , studio_name VARCHAR(50)  
  , producer_num INT  
  , PRIMARY KEY(title, year_made)  
);
```

```
DROP TABLE IF EXISTS MOVIE_STAR;  
CREATE TABLE MOVIE_STAR (  
  full_name VARCHAR(255)  
    PRIMARY KEY  
  , address VARCHAR(255)  
  , gener CHAR(1)  
  , birthdate DATE  
);
```

```
STARS_IN (  
  title:      string,  
  year_made: integer,  
  star_name: string  
)
```

```
MOVIE_EXECUTIVE (  
  full_name: string,  
  address:  string,  
  cert_num: integer,  
  net_worth: integer  
)
```

```
STUDIO (  
  legal_name:  string,  
  address:     string,  
  prez_cert_num: integer  
)
```



SQL I REVIEW EXERCISES...

```
DROP TABLE IF EXISTS MOVIE;  
CREATE TABLE MOVIE (  
    title VARCHAR(50)  
    , year_made INT  
    , length_mins INT  
    , genre VARCHAR(20)  
    , studio_name VARCHAR(50)  
    , producer_num INT  
    , PRIMARY KEY(title, year_made)  
);
```

```
DROP TABLE IF EXISTS MOVIE_STAR;  
CREATE TABLE MOVIE_STAR (  
    full_name VARCHAR(255)  
        PRIMARY KEY  
    , address VARCHAR(255)  
    , gener CHAR(1)  
    , birthdate DATE  
);
```

```
DROP TABLE IF EXISTS STARS_IN;  
CREATE TABLE STARS_IN (  
    title VARCHAR(50)  
    , year_made INT  
    , star_name VARCHAR(255)  
    , PRIMARY KEY(title  
        , year_made  
        , star_name)  
);
```

```
MOVIE_EXECUTIVE (  
    full_name: string,  
    address:   string,  
    cert_num: integer,  
    net_worth: integer  
)
```

```
STUDIO (  
    legal_name:   string,  
    address:      string,  
    prez_cert_num: integer  
)
```



SQL I REVIEW EXERCISES...

```
DROP TABLE IF EXISTS MOVIE;
CREATE TABLE MOVIE (
    title VARCHAR(50)
    , year_made INT
    , length_mins INT
    , genre VARCHAR(20)
    , studio_name VARCHAR(50)
    , producer_num INT
    , PRIMARY KEY(title, year_made)
);
```

```
DROP TABLE IF EXISTS MOVIE_STAR;
CREATE TABLE MOVIE_STAR (
    full_name VARCHAR(255)
        PRIMARY KEY
    , address VARCHAR(255)
    , gener CHAR(1)
    , birthdate DATE
);
```

```
DROP TABLE IF EXISTS STARS_IN;
CREATE TABLE STARS_IN (
    title VARCHAR(50)
    , year_made INT
    , star_name VARCHAR(255)
    , PRIMARY KEY(title
        , year_made
        , star_name)
);
```

```
DROP TABLE IF EXISTS
    MOVIE_EXECUTIVE;
CREATE TABLE
    MOVIE_EXECUTIVE (
        full_name VARCHAR(255)
    , address VARCHAR(255)
    , cert_num INT PRIMARY KEY
    , net_worth INT
    );
```

```
STUDIO (
    legal_name:    string,
    address:       string,
    prez_cert_num: integer
)
```



SQL I REVIEW EXERCISES...

```
DROP TABLE IF EXISTS MOVIE;  
CREATE TABLE MOVIE (  
    title VARCHAR(50)  
    , year_made INT  
    , length_mins INT  
    , genre VARCHAR(20)  
    , studio_name VARCHAR(50)  
    , producer_num INT  
    , PRIMARY KEY(title, year_made)  
);
```

```
DROP TABLE IF EXISTS MOVIE_STAR;  
CREATE TABLE MOVIE_STAR (  
    full_name VARCHAR(255)  
        PRIMARY KEY  
    , address VARCHAR(255)  
    , gener CHAR(1)  
    , birthdate DATE  
);
```

```
DROP TABLE IF EXISTS STARS_IN;  
CREATE TABLE STARS_IN (  
    title VARCHAR(50)  
    , year_made INT  
    , star_name VARCHAR(255)  
    , PRIMARY KEY(title  
        , year_made  
        , star_name)  
);
```

```
DROP TABLE IF EXISTS  
    MOVIE_EXECUTIVE;  
CREATE TABLE  
    MOVIE_EXECUTIVE (  
    full_name VARCHAR(255)  
    , address VARCHAR(255)  
    , cert_num INT PRIMARY KEY  
    , net_worth INT  
);
```

```
DROP TABLE IF EXISTS STUDIO;  
CREATE TABLE STUDIO (  
    legal_name VARCHAR(100)  
    , address VARCHAR(255)  
    , prez_cert_num INT PRIMARY KEY  
);
```



Maria DB Auto Increment

```
DROP TABLE IF EXISTS Location;  
CREATE TABLE Location (  
    LocationId INT  
        PRIMARY KEY AUTO_INCREMENT  
, LocationName VARCHAR(100)  
        UNIQUE  
) AUTO_INCREMENT = 101;
```

```
INSERT INTO Location (LocationName)  
VALUES ('Location101');
```

```
SELECT * FROM Location
```



Maria DB Auto Increment

```
DROP TABLE IF EXISTS Location;
CREATE TABLE Location (
  LocationId INT
    PRIMARY KEY AUTO_INCREMENT
, LocationName VARCHAR(100)
  UNIQUE
) AUTO_INCREMENT = 101;
```

```
INSERT INTO Location (LocationName)
VALUES ('Location101');
```

```
SELECT * FROM Location
```

Starts at 101 like we'd expect!

<input type="checkbox"/>	LocationId	LocationName
<input type="checkbox"/>	101	Location101

W

Maria DB Auto Increment

```
INSERT INTO Location (LocationName)
VALUES ('Location101');
```

<Oops... Error, cannot insert due to unique key constraint>

```
INSERT INTO Location (LocationName)
VALUES ('Location102');
```

```
SELECT * FROM Location
```

What do you expect will be the results?



Maria DB Auto Increment

```
INSERT INTO Location (LocationName)
VALUES ('Location101');
```

<Oops... Error, cannot insert due to unique key constraint>

```
INSERT INTO Location (LocationName)
VALUES ('Location102');
```

```
SELECT * FROM Location
```

LocationId	LocationName
101	Location101
103	Location102

W

Insert the following Tuples into Movie

title	year_made	length_mins	genre	studio_name	producer_num
Galaxy Quest	1999	104	comedy	DreamWorks	67890
Star Wars	1977	124	sciFi	Fox	12345
Wayne's World	1992	95	comedy	NULL	NULL



Insert the following Tuples into Movie

title	year_made	length_mins	genre	studio_name	producer_num
Galaxy Quest	1999	104	comedy	DreamWorks	67890
Star Wars	1977	124	sciFi	Fox	12345
Wayne's World	1992	95	comedy	NULL	NULL

INSERT INTO MOVIE VALUES

```
( 'Galaxy Quest' , 1999, 104, 'comedy', 'DreamWorks', 67890)
, ( 'Star Wars' , 1977, 104, 'sciFi' , 'Fox' , 12345)
, ( 'Wayne\'s World', 1992, 95, 'comedy', NULL , NULL);
```



Insert the following Tuples into Movie

title	year_made	length_mins	genre	studio_name	producer_num
The Matrix	1999	150	NULL	NULL	NULL



Insert the following Tuples into Movie

title	year_made	length_mins	genre	studio_name	producer_num
The Matrix	1999	150	NULL	NULL	NULL

```
INSERT INTO MOVIE (year_made, length_mins, title)
VALUES (1999, 150, 'The Matrix');
```





UNIVERSITY *of* WASHINGTON

JOINS



W

Implicit Joins vs Explicit Joins

```
SELECT *  
  FROM t1, t2  
 WHERE t1.b = t2.b
```

$$\sigma_{t1.b = t2.b}(t1 \times t2)$$

```
SELECT *  
  FROM t1  
 JOIN t2  
   ON t2.b = t2.b
```

$$t1 \bowtie_{t1.b=t2.b} t2$$


JOIN EXAMPLES

```
MOVIE (title, year_made, length_mins_genre, studio_name, producer_num)
MOVIE_STAR (name, address, gender, birthdate)
STARS_IN (title, year_made, star_name)
MOVIE_EXECUTIVE (full_name, address, cert_num, net_worth)
STUDIO (legal_name, address, prez_cert_num)
```

Write the following queries in SQL:

Who were the male stars in *Titanic*?

Which stars appeared in movies produced by MGM in 1995?

Who is the president of MGM studios?



JOIN EXAMPLES

Who were the male stars in *Titanic*?

```
SELECT MS.name  
  
FROM MOVIE_STAR MS  
JOIN STARS_IN SI  
    ON MS.name = SI.star_name  
WHERE SI.title = 'Titanic'  
      AND MS.gender = 'M'
```



JOIN EXAMPLES

Which stars appeared in movies produced by MGM in 1995?

```
SELECT DISTINCT SI.star_name
FROM STARS_IN SI
JOIN MOVIE M
  ON SI.title = M.title
  AND SI.year_made = M.year_made
WHERE M.studio_name = 'MGM'
      AND M.year_made = 1995
```



JOIN EXAMPLES

Who is the president of MGM studios?

```
SELECT ME.full_name  
  
FROM STUDIO S  
JOIN MOVIE_EXECUTIVE ME  
  ON S.prez_cert_num = ME.cert_num  
  
WHERE S.legal_name = 'MGM';
```





UNIVERSITY *of* WASHINGTON

Subqueries



W

What are Subqueries

- > Queries that are embedded into other queries are called “Subqueries”
- > Subqueries can sometimes run many levels deep
- > Use sparingly, can really decrease performance if misused



What can subqueries return?

- > A single “scalar” value:
 - (e.g. `SELECT COUNT(*) FROM MOVIE WHERE year = 1995`)
- > A relation of a single column, multiple rows
 - (e.g. `SELECT title FROM MOVIE WHERE year = 1995`)
- > A relation of multiple rows and columns
 - (e.g. `SELECT name, studio FROM MOVIE`)



Where are subqueries used?

> Where they are used depends on what they return

> **Scalar Return:**

– **SELECT CLAUSE**

```
SELECT
  MS.full_name
, (SELECT MIN(SI.year_made)
    FROM STARS_IN SI
    WHERE SI.star_name = MS.full_name)
  AS YEAR_OF_FIRST_MOVIE

FROM MOVIE_STAR MS
```



Where are subqueries used?

- > Where they are used depends on what they return
- > **Scalar Return:**
 - **WHERE CLAUSE**

```
SELECT ME1.*  
  FROM MOVIE_EXECUTIVE ME1  
 WHERE ME1.net_worth >=  
       (SELECT AVG(ME2.net_worth)  
        FROM MOVIE_EXECUTIVE ME2)
```



Where are subqueries used?

> Where they are used depends on what they return

> Single Column Return

```
SELECT
    E.EmployeeName
    , E.JobTitle
FROM Employee E
WHERE EmployeeId IN
    (SELECT EmployeeId
     FROM Transcripts T
     WHERE T.Status = 'Complete')
```



Where are subqueries used?

- > Where they are used depends on what they return
- > Multiple Rows, Multiple Columns
 - FROM / JOIN

```
SELECT
    T.COL1, T.COL2
FROM
    (SELECT
        /* SOME COMPLICATED QUERY */
        FROM
            /* A BUNCH OF LOGIC */
        ) T

JOIN MOVIE E
ON T.Col1 = E.MovieName
```



Operators that Deal with Relations

> **EXISTS**

The subquery returns at least one record

> **IN**

Similar to a long “or” statement

– Age in (Subquery)... Age = 1 or Age = 2, or Age = 5)

> **ANY**

The condition holds for a single value

– Age > ANY (Subquery)

> **ALL**

The condition holds for all values

– Age > ALL (Subquery)



More Examples (IN)

```
1) SELECT name
2) FROM MovieExec
3) WHERE cert# IN
4)     (SELECT producerC#
5)     FROM Movies
6)     WHERE (title, year) IN
7)     (SELECT movieTitle, movieYear
8)     FROM StarsIn
9)     WHERE starName = 'Harrison Ford'
    )
);
```

More Examples (EXISTS)

List the full cast from movies who have at least one movie executive starring in the movie.

```
SELECT *
FROM STARS_IN SI

WHERE EXISTS
    (SELECT ME.full_name
     FROM MOVIE_EXECUTIVE ME
     JOIN STARS_IN SI2
       ON ME.full_name = SI2.full_name
     WHERE SI.title      = SI2.title
      AND SI.year_made = SI2.year_made)
```



More Examples (ANY | ALL)

Find the brand and wattages of lightbulbs that are not the brightest. Second, find the ones that are

```
SELECT  
    B.Brand  
    , B.Wattage as "NotTheBrightestBulbInTheBox"
```

```
FROM LightBulbs B
```

```
WHERE B.Wattage < ANY (SELECT B2.Wattage  
                        FROM LightBulbs B2)
```

```
SELECT  
    B.Brand  
    , B.Wattage as "TheBrightestBulbInTheBox"
```

```
FROM LightBulbs B
```

```
WHERE B.Wattage >= ALL (SELECT B2.Wattage  
                        FROM LightBulbs B2)
```



Correlated Queries

Some subqueries can need to be executed many times, essentially once for each tuple returned by the outer query.

Queries that contain these types of subqueries are called **Correlated Queries**.

```
SELECT
    MS.full_name
, (SELECT MIN(SI.year_made)
    FROM STARS_IN SI
    WHERE SI.star_name = MS.full_name)
  AS YEAR_OF_FIRST_MOVIE

FROM MOVIE_STAR MS
```



Correlated Queries

Can sometimes be rewritten, sometimes not. Rarely should be used on queries that return significantly large result sets.

```
SELECT
  MS.full_name
, (SELECT MIN(SI.year_made)
   FROM STARS_IN SI
   WHERE SI.star_name = MS.full_name)
  AS YEAR_OF_FIRST_MOVIE
FROM MOVIE_STAR MS
```

```
SELECT MS.full_name
      , MIN(SI.year_made)
      AS "YEAR_OF_FIRST_MOVIE"
FROM MOVIE_STAR MS
JOIN STARS_IN SI
  ON MS.full_name = SI.star_name
GROUP BY MS.full_name;
```





UNIVERSITY *of* WASHINGTON

OUTER JOINS AND CROSS JOINS



W

Outer Joins

- > Combine two tables on a given condition.
- > Dangling tuples will remain in the result.
 - From which table depends on the type of outer join
 - > FULL: Dangling Tuples from both tables remain
 - > LEFT: Dangling Tuples from the LEFT table remain
 - > RIGHT: Dangling Tuples from the RIGHT table remain



Outer Joins

- > The "LEFT" table is to the "LEFT" of the JOIN
- > THE "RIGHT" table is to the "RIGHT" of the JOIN

```
SELECT *  
FROM left_table _____ OUTER JOIN right_table
```



Outer Joins Examples: LEFT OUTER

T1	
A	B
1	2
2	3
3	2
4	6

T2	
B	C
1	Shake It All About
2	Left Hand In
3	Left Hand Out
4	What's It All About?

```
SELECT *  
FROM T1  
LEFT OUTER JOIN T2  
  ON T1.B = T2.B
```



Outer Joins Examples: LEFT OUTER

T1	
A	B
1	2
2	3
3	2
4	6

T2	
B	C
1	Shake It All About
2	Left Hand In
3	Left Hand Out
4	What's It All About?

```
SELECT *  
FROM T1  
LEFT OUTER JOIN T2  
  ON T1.B = T2.B
```

A	T1.B	T2.B	C
1	2	2	Left Hand In
2	2	2	Left Hand Out
3	2	2	Left Hand In
4	6	NULL	NULL



Outer Joins Examples: RIGHT OUTER

T1	
A	B
1	2
2	3
3	2
4	6

T2	
B	C
1	Shake It All About
2	Left Hand In
3	Left Hand Out
4	What's It All About?

```
SELECT *  
FROM T1  
RIGHT OUTER JOIN T2  
  ON T1.B = T2.B
```



Outer Joins Examples: RIGHT OUTER

T1	
A	B
1	2
2	3
3	2
4	6

T2	
B	C
1	Shake It All About
2	Left Hand In
3	Left Hand Out
4	What's It All About?

```
SELECT *  
FROM T1  
RIGHT OUTER JOIN T2  
ON T1.B = T2.B
```

A	T1.B	T2.B	C
NULL	NULL	1	Shake it All About
1	2	2	Left Hand In
3	2	2	Left Hand In
2	3	3	Left Hand Out
NULL	NULL	4	What's It All About



Outer Joins Examples: FULL OUTER

T1	
A	B
1	2
2	3
3	2
4	6

T2	
B	C
1	Shake It All About
2	Left Hand In
3	Left Hand Out
4	What's It All About?

```
SELECT *  
FROM T1  
FULL OUTER JOIN T2  
  ON T1.B = T2.B
```



Outer Joins Examples: FULL OUTER

T1	
A	B
1	2
2	3
3	2
4	6

T2	
B	C
1	Shake It All About
2	Left Hand In
3	Left Hand Out
4	What's It All About?

```
SELECT *  
FROM T1  
FULL OUTER JOIN T2  
  ON T1.B = T2.B
```

A	T1.B	T2.B	C
1	2	2	Left Hand In
2	3	3	Left Hand Out
3	2	2	Left Hand In
4	6	NULL	NULL
NULL	NULL	1	Shake it All About
NULL	NULL	4	What's It All About?



Cross Joins (Cartesian Product)

- > Every tuple of the left table is paired with every tuple of the right table. $|R| \times |S| = |R \times S|$

R		S	
A	B	C	D
1	2	5	6
3	4	7	8

A	B	C	D
1	2	5	6
1	2	7	8
3	4	5	6
3	4	7	8



Cross Join Syntax

```
SELECT *  
FROM R  
CROSS JOIN S
```

```
SELECT *  
FROM R,S
```



Fun with Cross Joins

```
CREATE TABLE B (C INT);  
INSERT INTO B VALUES (0), (1);
```

C
0
1



Fun with Cross Joins

```
CREATE TABLE B (C INT);  
INSERT INTO B VALUES (0), (1);
```

C
0
1

WHAT WILL B x B x B x B LOOK LIKE?



Fun with Cross Joins

```
SELECT *  
FROM B B1, B B2, B B3, B B4  
ORDER BY B1.C, B2.C, B3.C, B4.C
```



Fun with Cross Joins

```
SELECT *  
FROM B B1, B B2, B B3, B B4  
ORDER BY B1.C, B2.C, B3.C, B4.C
```

c	c	c	c
0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
1	0	0	1
1	0	1	0
1	0	1	1
1	1	0	0
1	1	0	1
1	1	1	0
1	1	1	1

W

Use Case: For Each Day

Provide a list that includes every day, every shift, how many cogs were made.

DATE_DIM

SHORT_DATE	DAY_OF_WEEK
1/1/17	SUN
1/2/17	MON
1/3/17	TUE
1/4/17	WED
...	
12/31/99	FRI

COGS_BUILT

DATE	COGS	SHIFT	SHIFTS
1/1/17	4	AM	AM
1/2/17	5	PM	PM
1/4/17	10	AM	
1/4/17	20	PM	

DATE	SHIFT	COGS
1/1/17	AM	4
1/1/17	PM	0
1/2/17	AM	0
1/2/17	PM	5
1/3/17	AM	0
1/3/17	PM	0
1/4/17	AM	10
1/4/17	PM	20

W

Use Case: For Each Day

START WITH A CROSS JOIN TO BUILD YOUR SCAFFOLD

```
(      SELECT DD.SHORT_DATE, S.SHIFT
      FROM DATE_DIM DD
      CROSS JOIN SHIFTS S
      WHERE DD.SHORT_DATE >= '2017-01-01'
            AND DD.SHORT_DATE < '2017-01-05'
) AS SCAFFOLD
```

SHORT_DATE	SHIFTS
1/1/17	AM
1/1/17	PM
1/2/17	AM
1/2/17	PM
1/3/17	AM
1/3/17	PM
1/4/17	AM
1/4/17	PM



Use Case: For Each Day

THEN OUTER JOIN COGS_BUILT TO THE SCAFFOLD TO GET THE COG COUNTS

```
SELECT
    SCAFFOLD.SHORT_DATE AS "DATE"
    , SCAFFOLD.SHIFT
    , C.COGS
FROM COGS_BUILT C
RIGHT OUTER JOIN
    (
        SELECT DD.SHORT_DATE, S.SHIFT
        FROM DATE_DIM DD
        CROSS JOIN SHIFTS S
        WHERE DD.SHORT_DATE >= '2017-01-01'
        AND DD.SHORT_DATE < '2017-01-05'
    ) AS SCAFFOLD
ON C.DATE = SCAFFOLD.SHORT_DATE
AND C.SHIFT = SCAFFOLD.SHIFT
```

DATE	SHIFT	COGS
1/1/17	AM	4
1/1/17	PM	NULL
1/2/17	AM	NULL
1/2/17	PM	5
1/3/17	AM	NULL
1/3/17	PM	NULL
1/4/17	AM	10
1/4/17	PM	20



Use Case: For Each Day

LASTLY WE HAVE TO HANDLE NULLS! IN MARIADB, USE IFNULL(value,value_if_null)

```
SELECT
    SCAFFOLD.SHORT_DATE AS "DATE"
    , SCAFFOLD.SHIFT
    , IFNULL(C.COGS, 0) AS "COGS"
FROM COGS_BUILT C
RIGHT OUTER JOIN
    (
        SELECT DD.SHORT_DATE, S.SHIFT
        FROM DATE_DIM DD
        CROSS JOIN SHIFTS S
        WHERE DD.SHORT_DATE >= '2017-01-01'
        AND DD.SHORT_DATE < '2017-01-05'
    ) AS SCAFFOLD
ON C.DATE = SCAFFOLD.SHORT_DATE
AND C.SHIFT = SCAFFOLD.SHIFT
```

DATE	SHIFT	COGS
1/1/17	AM	4
1/1/17	PM	0
1/2/17	AM	0
1/2/17	PM	5
1/3/17	AM	0
1/3/17	PM	0
1/4/17	AM	10
1/4/17	PM	20

W



UNIVERSITY *of* WASHINGTON

Natural Joins



W

Natural Joins. Possible in SQL

Matches on columns with identical names.

Depends on a well formed schema.

```
SELECT *  
      FROM Employee  
NATURAL JOIN TrainingStatus
```



A look at all of the joins...

FROM R, S

FROM R NATURAL JOIN S

FROM R NATURAL [LEFT|RIGHT|FULL] OUTER JOIN S

FROM R CROSS JOIN S

FROM R JOIN S ON...

FROM R [LEFT|RIGHT|FULL] OUTER JOIN S ON...





UNIVERSITY *of* WASHINGTON

Grouping & Aggregation



W

Aggregation

SUM, AVG, MIN, MAX, COUNT

SELECT COUNT(*) FROM MOVIE

SELECT COUNT(title) FROM MOVIE

SELECT COUNT(DISTINCT star_name) FROM STARS_IN



Aggregation With “Group By”

```
SELECT AVG(length_mins) as "avg_length"  
      , COUNT(*)          as "movies_made"  
FROM MOVIE  
GROUP BY studio_name;
```

```
SELECT year_made  
      , count(*) as "movies_made"  
FROM MOVIE  
GROUP BY year_made;
```



Aggregation With “Group By”

```
SELECT full_name  
FROM STARS_IN  
GROUP BY full_name
```

EXACTLY THE SAME RESULT AS

```
SELECT DISTINCT full_name  
FROM STARS_IN
```



Aggregate Attributes & Grouping Attributes

Why will the following query fail to return the correct results?

```
SELECT studio, year_made, count(*) as release_count  
FROM MOVIE  
GROUP BY year_made
```

What do we have to change to get it to run?



Aggregate Attributes & Grouping Attributes

Why will the following query fail to return the correct results?

```
SELECT studio, year_made, count(*) as release_count  
FROM MOVIE  
GROUP BY year_made
```

What do we have to change to get it to run?

```
SELECT studio, year_made, count(*) as release_count  
FROM MOVIE  
GROUP BY year_made, studio
```



NULLS in Aggregation

- > NULL is (in general) ignored by aggregation
 - $\text{AVG}(1, \text{null}, 2, 3) = (1+2+3)/3 = 2$
- > `COUNT(column)` returns the number of non NULL values in that column
- > `COUNT(*)` is the number of tuples / rows regardless of nulls
- > READ DOCUMENTATION OF YOUR DBMS
 - Special cases exist when entire column or entire row is null
 - > MariaDB: `count(null) = 0`, `sum(null) = null`



HAVING

(Condition applied after Grouping)

Having is a “WHERE” condition that occurs AFTER aggregation takes place.

Example: get movies that cast more than 2 movie stars

```
SELECT title, year_made  
FROM STARS_IN  
  
GROUP BY title, year_made  
  
HAVING COUNT(*) > 2
```



HAVING

(Condition applied after Grouping)

Having is a “WHERE” condition that occurs AFTER aggregation takes place.

Example: get movies that cast more than 2 movie stars

```
SELECT title, year_made
FROM STARS_IN

GROUP BY title, year_made

HAVING COUNT(*) > 2
```

```
SELECT T.title, T.year_made
FROM
    ( SELECT title
      , year_made
      , count(*) as movie_count
    FROM STARS_IN SI2
    GROUP BY title, year_made
    ) as T

WHERE T.movie_count > 2
```





Additional Examples

Get movies from Fox that cast more than 2 movie stars

W

Additional Examples

Get movies from Fox that cast more than 2 movie stars

```
SELECT SI.title, SI.year_made  
  
      FROM STARS_IN SI  
     JOIN MOVIE M  
       ON SI.title = M.title  
       AND SI.year_made = M.year_made  
  
     WHERE M.studio_name = 'Fox'  
  
    GROUP BY SI.title, SI.year_made  
  
   HAVING COUNT(*) > 2
```





Additional Examples

Get the average length of movies from each studio that has produced at least 3 movies.



Additional Examples

Get the average length of movies from each studio that has produced at least 3 movies.

```
SELECT studio_name
       , avg(length_mins) as avg_len
  FROM MOVIE
 GROUP BY studio_name
  HAVING count(*) >= 3
```



Order By

- > The last statement executed (typically)
- > Often not allowed in subqueries (not needed)
- > Ordered columns do NOT need to be in projection
- > Can include Aggregated Columns

ORDER BY colname [ASC|DESC], colname [ASC|DESC]...



Order By Examples

List movie stars in alphabetical order

```
SELECT full_name, address  
FROM MOVIE_STAR  
ORDER BY full_name ASC
```



Order By Examples

List movie stars in order of how many movies they've starred in (most on top, ties sort in alphabetical order)

```
SELECT full_name, count(*) as movie_count  
FROM STARS_IN  
GROUP BY full_name  
ORDER BY count(*) DESC, full_name ASC
```





UNIVERSITY *of* WASHINGTON

Unions, Intersections, and Difference



W

Set Operations

- SQL provides corresponding operators that apply to the results of queries, provided those queries produce relations with the same list of attributes and attribute types.
- The keywords used are UNION, INTERSECT, and EXCEPT for \cup , \cap , and $-$, respectively.
- There is a bag union called “UNION ALL” that will not remove duplicates but rather stick one relation on top of the other.
- *Not all databases support these keywords.*



Set Operations

MovieStar (name, address, gender, birthdate)

MovieExec (name, address, cert#, netWorth)

Movies (title, year, length, genre, studioName, producerC#)


StarsIn (movieTitle, movieYear, starName)

```
SELECT name, address
  FROM MovieStar
 WHERE gender = 'F'
INTERSECT
SELECT name, address
  FROM MovieExec
 WHERE netWorth > 10000000;
```

```
SELECT name, address
  FROM MovieStar
EXCEPT
SELECT name, address
  FROM MovieExec;
```

```
SELECT title, year
  FROM Movies
UNION
SELECT movieTitle AS title
      , movieYear  AS year
  FROM StarsIn;
```





UNIVERSITY *of* WASHINGTON

MANIPULATING DATA THAT EXISTS



W



To Modify Data

- > INSERT**
- > UPDATE**
- > DELETE**

W

INSERT

**INSERT INTO table (column_list) VALUES
(value_list) , (value_list)**

**INSERT INTO table (column_list)
SELECT ...**

**To use select statement, you must have matching
tuple sizes and adhere to all constraints just as if you
were using VALUES (xxx),(xxx)**



DELETE

> DELETE FROM table WHERE condition...

```
DELETE FROM StarsIn  
WHERE movie_title = 'Star Wars';
```

WARNING:

REPLACE **DELETE FROM** WITH **SELECT * FROM** AS A SANITY CHECK BEFORE EXECUTING YOUR DELETE STATEMENT





DELETE

> DELETING ALL DATA FROM A TABLE:

```
DELETE FROM tablename;
```



UPDATE

UPDATE table SET <value assignment> WHERE <condition>

```
UPDATE MOVIE SET title = 'Star Wars: A New Hope'
              WHERE title = 'Star Wars'
              AND year = '1977'
```

```
UPDATE MOVIE_EXECUTIVE
      SET full_name = CONCAT('Pres. ', full_name)
      WHERE LEFT(full_name,6) != 'Pres. '
```

